

Convolutional recurrent neural networks for music classification

Published in: 2017 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)

7107018026

劉俊廷

Contents

CNN and CRNN for music classification

Introductions to the models

Experiments

References


CNN and CRNN for music classification

- CNN
 - CNNs assume features that are in different levels of hierarchy and can be extracted by convolutional kernels.
 - CNN for music classification:
 - music tagging: AUC:0.88~0.89
 - genre classification: accuracy:66~69%
 - user-item latent feature prediction for recommendation: AUC:0.77192


CNN and CRNN for music classification

- CRNN
 - CRNN
 - Modified CNN by replacing the last convolutional layers with a RNN.
 - Adopting an RNN for aggregating the features enables the networks to take the global structure into account while local features are extracted by the remaining convolutional layers.
 - CRNN for music classification:
 - Music transcription: Accuracy:72%~76%
 - Document classification:Accuracy:65%~67%
 - Image classification:Accuracy:51.4%


Introductions to the models




(a) k1c2



(b) k2c1



(c) k2c2



(d) CRNN

Sigmoid functions are used as activation at output nodes because music tagging is a *multi-label* classification task.

Experiments

- Dataset
 - the Million Song Dataset with *last.fm* tags.
 - Predict the top-50 tag, which includes **genres** (e.g., *rock*), **moods** (e.g., *happy*), **instruments** (e.g., *guitar*), and **eras** (*60s–00s*).
 - 214285 (201680 for training and 12605 for validation).
- Input: Resulting in an input shape of **96×1366** (mel-frequency band×time frame).
- We use **ADAM** for learn and **binary cross-entropy** as a loss function.
- Computes **AUC-ROC** for compare four models
- The result of this paper, the AUC of CRNN can arrive 0.85~0.87.

Thanks for listening

References

- Keunwoo Choi, George Fazekas, Mark Sandler, "Automatic tagging using deep convolutional neural networks"
- Paulo Chiliguano, Gyorgy Fazekas, "Hybrid music recommender using content-based and social information"
- Aaron Van den Oord, Sander Dieleman, Benjamin Schrauwen, "Deep content-based music recommendation"
- Duyu Tang, Bing Qin, Ting Liu, "Document modeling with gated recurrent neural network for sentiment classification"
- Zhen Zuo, Bing Shuai, Gang Wang, Xiao Liu, Xingxing Wang, Bing Wang, Yushi Chen, "Convolutional recurrent neural networks: Learning spatial dependencies for image representation",
- Siddharth Sigtia, Emmanouil Benetos, Simon Dixon, "An end-to-end neural network for polyphonic piano music transcription"
- Diederik P. Kingma, Jimmy Ba, Adam: A method for stochastic optimization